

AB_Testing

Reinp

2020-05-05

```
## Set Chunk requirements
knitr::opts_chunk$set(echo = TRUE, message = FALSE, warning = FALSE)
#echo=FALSE indicates that the code will not be shown in the final document
#(though any results/output would still be displayed).
#include=FALSE to have the chunk evaluated, but neither the code nor its output displayed
# warning=FALSE and message=FALSE suppress any R warnings or messages from being included
#in the final document

#Import relevant packages
library(stats)
library(psych)
library(ggplot2)

#loading the excel dataset with two sheets
library("readxl")
setwd('E:/Documents/Reinp/GitHub Respositories/AB_Testing-with-RStudio')
ds_sd<-read_excel('School data.xlsx', sheet = "Data")
View(ds_sd)
attach(ds_sd)

ds_sd1<-read_excel('School data.xlsx', sheet = "District names")
View(ds_sd1)
attach(ds_sd1)

#learn more about the dataset
help(ds_sd)

## No documentation for 'ds_sd' in specified packages and libraries:
## you could try '??ds_sd'

??ds_sd
str(ds_sd)

## Classes 'tbl_df', 'tbl' and 'data.frame':   6984 obs. of  6 variables:
## $ district_id      : num  78 78 78 78 78 78 78 78 78 78 ...
## $ attended_male    : num  208 337 215 200 128 185 216 180 183 205 ...
## $ attended_female  : num  205 281 162 167 137 140 168 130 159 161 ...
## $ enrolled_male_students : num  769 859 695 850 609 556 477 512 604 225 ...
## $ enrolled_female_students: num  620 685 585 701 519 442 342 436 469 176 ...
## $ treatment        : num  1 0 0 0 1 0 0 1 1 1 ...

class(ds_sd)

## [1] "tbl_df"      "tbl"        "data.frame"
```

```

typeof(ds_sd)

## [1] "list"
length(ds_sd)

## [1] 6
names(ds_sd) #display variable names

## [1] "district_id"          "attended_male"
## [3] "attended_female"      "enrolled_male_students"
## [5] "enrolled_female_students" "treatment"

#attributes(ds_sd) #names(ds_sd), class(ds_sd), row.names(ds_sd)

#distribution of input variables
table(ds_sd$district_id)

##
##      3      5      6     16     17     22     39     42     43     44     47     50     57     58     62     65     67     68     77     78
## 253 279 256 302 218 227 274 225 248 237 222 241 294 251 239 290 260 243 260 270
##   79   80   93   97  104  107  108  200
## 319 288 245 247 233 254 295   14

table(ds_sd1$"DISTRICT ID")

##
##      1      2      3      5      6     16     17     22     39     42     43     44     47     50     57     58     62     65     67     68
##      1      1      1      1      1      1      1      1      1      1      1      1      1      1      1      1      1      1      1
##  77  78  79  80  93  97 104 107 108 112
##      1      1      1      1      1      1      1      1      1      1

#unique values per column
unique(ds_sd$district_id)

## [1] 78 79 3 80 39 5 6 16 17 22 42 43 93 44 97 47 50 57 58
## [20] 62 104 65 107 108 67 68 77 200

unique(ds_sd1$"DISTRICT ID")

## [1] 78 79 3 80 39 5 6 16 17 22 42 43 93 44 97 47 50 57 58
## [20] 62 104 65 107 108 67 68 77 112 1 2

#district_id 200 in (ds_sd dataframe) is not in the (ds_sd1 dataframe)
#DISTRICT ID 1 in (ds_sd1 dataframe) is not in the (ds_sd dataframe)
#DISTRICT ID 2 in (ds_sd1 dataframe) is not in the (ds_sd dataframe)
#DISTRICT ID 112 in (ds_sd1 dataframe) is not in the (ds_sd dataframe)

ds_sdmerge <- merge(ds_sd, ds_sd1, by.x = "district_id", by.y = "DISTRICT ID") #N/A.
#This is default #Keep rows where there's a match in both #innerjoin
View(ds_sdmerge)

ds_sdmerge1 <- merge(ds_sd, ds_sd1, by.x = "district_id", by.y = "DISTRICT ID", all.x = TRUE)
#Keep all rows from x, regardless of match in y #leftJoin even if there's no match in y
View(ds_sdmerge1)

ds_sdmerge2 <- merge(ds_sd, ds_sd1, by.x = "district_id", by.y = "DISTRICT ID", all.y = TRUE)
#Keep all rows from y, regardless of match in x #Rightjoin even if there's no match in x

```

```

View(ds_sdmerge2)

ds_sdmerge3 <- merge(ds_sd, ds_sd1, by.x = "district_id", by.y = "DISTRICT ID", all= TRUE)
#Keep all rows from x AND from y #Outerjoin
View(ds_sdmerge3)

#We select default merge for the final data. It keeps rows where there's a match in both.
#we get a total of 6970 schools

#add total_enrolled column by computing the total number of student enrolled and total number
#of attended in each school
#Get the total number of student enrolled in each school by adding enrolled_male_students
#and enrolled_female_students columns

#Adding by equation
ds_sdmerge$total_enrolled=ds_sdmerge$enrolled_male_students + ds_sdmerge$enrolled_female_students
ds_sdmerge$total_attended=ds_sdmerge$attended_male + ds_sdmerge$attended_female

#transform() function
ds_sdmerge1 <- transform(ds_sdmerge1, total_enrolled = enrolled_male_students +
                        enrolled_female_students)
ds_sdmerge1 <- transform(ds_sdmerge1, total_attended = attended_male + attended_female)

#apply() function
ds_sdmerge2$total_enrolled <- apply(ds_sdmerge2[,c('enrolled_male_students',
            'enrolled_female_students')], 1, function(x) sum(x))
ds_sdmerge2$total_attended <- apply(ds_sdmerge2[,c('attended_male', 'attended_female')], 1,
            function(x) sum(x))

#tidyverse's dplyr
library(dplyr)
ds_sdmerge3 <- mutate(ds_sdmerge3, total_enrolled = enrolled_male_students +
                        enrolled_female_students)
ds_sdmerge3 <- mutate(ds_sdmerge3, total_attended = attended_male + attended_female)

#summary statistics
summary(ds_sdmerge) #summarizes the dataset

```

```

##   district_id   attended_male   attended_female enrolled_male_students
##   Min.    : 3.00   Min.    : 0.0   Min.    : 0.0   Min.    : -999.0
##   1st Qu.: 39.00   1st Qu.:190.0   1st Qu.:161.0   1st Qu.: 321.0
##   Median : 58.00   Median :247.0   Median :207.0   Median : 479.0
##   Mean   : 57.27   Mean   :253.3   Mean   :211.6   Mean   : 531.9
##   3rd Qu.: 79.00   3rd Qu.:308.0   3rd Qu.:257.0   3rd Qu.: 730.0
##   Max.    :108.00   Max.    :597.0   Max.    :483.0   Max.    :1194.0
## enrolled_female_students treatment District Name total_enrolled
##   Min.    : -999.0   Min.    :0.0000   Length:6970   Min.    : -1998.0
##   1st Qu.: 268.0   1st Qu.:0.0000   Class :character   1st Qu.: 592.0
##   Median : 394.0   Median :0.0000   Mode  :character   Median : 870.5
##   Mean   : 438.7   Mean   :0.4973           Mean   : 970.7
##   3rd Qu.: 600.8   3rd Qu.:1.0000           3rd Qu.:1331.8
##   Max.    : 991.0   Max.    :1.0000           Max.    :1999.0
## total_attended
##   Min.    : 0.0

```

```
## 1st Qu.: 353.0
## Median : 456.0
## Mean : 464.9
## 3rd Qu.: 564.0
## Max. :1045.0
```

```
describe(ds_sdmerge)
```

```
##          vars      n  mean      sd median trimmed   mad   min
## district_id      1 6970  57.27  31.22   58.0   57.68  29.65    3
## attended_male     2 6970 253.31  85.78  247.0  249.84  87.47    0
## attended_female   3 6970 211.62  70.68  207.0  208.50  71.16    0
## enrolled_male_students 4 6970 531.93 266.40 479.0 518.04 276.50 -999
## enrolled_female_students 5 6970 438.73 222.13 394.0 427.54 223.87 -999
## treatment         6 6970   0.50   0.50   0.0   0.50   0.00    0
## District Name*    7 6970   NaN    NA    NA    NaN    NA   Inf
## total_enrolled     8 6970 970.66 485.92 870.5 946.60 496.67 -1998
## total_attended     9 6970 464.93 153.75 456.0 458.95 157.16    0
##          max range  skew kurtosis   se
## district_id    108  105 -0.15   -0.92 0.37
## attended_male   597  597  0.39    0.09 1.03
## attended_female 483  483  0.42    0.13 0.85
## enrolled_male_students 1194 2193 -0.12    2.20 3.19
## enrolled_female_students 991 1990 -0.38    4.23 2.66
## treatment        1    1  0.01   -2.00 0.01
## District Name*  -Inf -Inf    NA    NA  NA
## total_enrolled  1999 3997 -0.26    3.07 5.82
## total_attended  1045 1045  0.36    0.03 1.84
```

#Create the school_id variable by first sorting the data within each district by the total number of enrollees per school. Let the ID be 1 for the school within each district with the highest number of enrolled students, 2 for the second highest and so on.....

```
ds_sdmerge <- arrange(ds_sdmerge,district_id,desc(total_enrolled))
```

```
ds_sdmerger <- ds_sdmerge %>%
  group_by(district_id) %>%
  mutate(school_id = rank(desc(total_enrolled), ties.method = "first"))
```

```
View(ds_sdmerger)
```

```
head(ds_sdmerger)
```

```
## # A tibble: 6 x 10
## # Groups:   district_id [1]
##   district_id attended_male attended_female enrolled_male_s~ enrolled_female~
##           <dbl>         <dbl>         <dbl>         <dbl>         <dbl>
## 1             3             212             177             1020             959
## 2             3             335             331             985             962
## 3             3             261             186             1139             807
## 4             3             208             166             1014             920
## 5             3             238             184             1050             881
## 6             3             112              75             1016             908
## # ... with 5 more variables: treatment <dbl>, `District Name` <chr>,
## #   total_enrolled <dbl>, total_attended <dbl>, school_id <int>
```

```
tail(ds_sdmerger)
```

```
## # A tibble: 6 x 10
## # Groups:   district_id [1]
##   district_id attended_male attended_female enrolled_male_s~ enrolled_female~
##         <dbl>         <dbl>         <dbl>         <dbl>         <dbl>
## 1         108         118           97          136          110
## 2         108          73           55          126          106
## 3         108          66           64          112          108
## 4         108         242          230         -999         -999
## 5         108         156          162         -999         -999
## 6         108         239          222         -999         -999
## # ... with 5 more variables: treatment <dbl>, `District Name` <chr>,
## #   total_enrolled <dbl>, total_attended <dbl>, school_id <int>
```

#Check the numeric variables for outliers.


```
describe(ds_sdmerger)
```

```
##           vars    n  mean    sd median trimmed   mad   min
## district_id      1 6970  57.27  31.22   58.0   57.68  29.65    3
## attended_male     2 6970 253.31  85.78  247.0  249.84  87.47    0
## attended_female   3 6970 211.62  70.68  207.0  208.50  71.16    0
## enrolled_male_students 4 6970 531.93 266.40  479.0  518.04 276.50 -999
## enrolled_female_students 5 6970 438.73 222.13  394.0  427.54 223.87 -999
## treatment         6 6970   0.50   0.50    0.0    0.50   0.00    0
## District Name*    7 6970   NaN    NA    NA    NaN    NA    Inf
## total_enrolled    8 6970 970.66 485.92  870.5  946.60 496.67 -1998
## total_attended    9 6970 464.93 153.75  456.0  458.95 157.16    0
## school_id        10 6970 130.90  76.85  130.0  129.61  96.37    1
##
##           max range  skew kurtosis   se
## district_id    108   105 -0.15   -0.92 0.37
## attended_male   597   597  0.39    0.09 1.03
## attended_female  483   483  0.42    0.13 0.85
## enrolled_male_students 1194 2193 -0.12    2.20 3.19
## enrolled_female_students 991 1990 -0.38    4.23 2.66
## treatment        1     1  0.01   -2.00 0.01
## District Name*  -Inf  -Inf    NA    NA    NA
## total_enrolled   1999 3997 -0.26    3.07 5.82
## total_attended   1045 1045  0.36    0.03 1.84
## school_id        319   318  0.11   -1.04 0.92
```

We use boxplot to visualize for any outliers

```
boxplot(ds_sdmerger[, c("attended_male", "attended_female", "enrolled_male_students",
                        "enrolled_female_students", "total_enrolled", "total_attended")], main="boxplots",
        xlab="variables",
        ylab="number of students",
        col="orange",
        border="brown", las = 2, cex.axis = 0.6, col.axis = 'blue', col.lab = 'red')
```

boxplots



*#From the boxplot above, outliers are present in form of negative count of students.
 #There are negative values in enrolled_male_students and in enrolled_female_students leading
 #to a negative total_enrolled*

#cleaning data from the outliers

```
ds_sdmerger1 <- ds_sdmerger[(ds_sdmerger[,8]>0),]
View(ds_sdmerger1)
describe(ds_sdmerger1)
```

| ## | vars | n | mean | sd | median | trimmed | mad | min | max |
|-----------------------------|------|------|--------|--------|--------|---------|--------|-----|------|
| ## district_id | 1 | 6948 | 57.27 | 31.22 | 58.0 | 57.67 | 29.65 | 3 | 108 |
| ## attended_male | 2 | 6948 | 253.42 | 85.76 | 247.0 | 249.93 | 87.47 | 0 | 597 |
| ## attended_female | 3 | 6948 | 211.71 | 70.67 | 207.0 | 208.57 | 71.16 | 0 | 483 |
| ## enrolled_male_students | 4 | 6948 | 536.78 | 252.49 | 481.0 | 519.18 | 277.25 | 106 | 1194 |
| ## enrolled_female_students | 5 | 6948 | 443.28 | 207.20 | 395.0 | 428.48 | 223.87 | 88 | 991 |
| ## treatment | 6 | 6948 | 0.50 | 0.50 | 0.0 | 0.50 | 0.00 | 0 | 1 |
| ## District Name* | 7 | 6948 | NaN | NA | NA | NaN | NA | Inf | -Inf |
| ## total_enrolled | 8 | 6948 | 980.06 | 457.02 | 873.5 | 948.68 | 495.93 | 204 | 1999 |
| ## total_attended | 9 | 6948 | 465.13 | 153.72 | 456.0 | 459.11 | 157.16 | 0 | 1045 |
| ## school_id | 10 | 6948 | 130.49 | 76.62 | 129.0 | 129.21 | 94.89 | 1 | 319 |

| ## | range | skew | kurtosis | se |
|-----------------------------|-------|-------|----------|------|
| ## district_id | 105 | -0.15 | -0.92 | 0.37 |
| ## attended_male | 597 | 0.40 | 0.08 | 1.03 |
| ## attended_female | 483 | 0.42 | 0.12 | 0.85 |
| ## enrolled_male_students | 1088 | 0.50 | -0.86 | 3.03 |
| ## enrolled_female_students | 903 | 0.52 | -0.81 | 2.49 |
| ## treatment | 1 | 0.01 | -2.00 | 0.01 |

```
## District Name*          -Inf      NA      NA      NA
## total_enrolled          1795  0.49   -0.90  5.48
## total_attended          1045  0.37    0.03  1.84
## school_id               318  0.11   -1.04  0.92
```

```
#Label values for the treatment variable appropriately (1 = Treatment, 0 = Control)
ds_sdmerger1$treatment_type <- factor(ds_sdmerger1$treatment, levels = c(0,1),
                                     labels = c("control", "treatment"))
```

```
View(ds_sdmerger1)
```

```
which(is.na(ds_sdmerger1$total_attended)) #check for missing values
```

```
## integer(0)
```

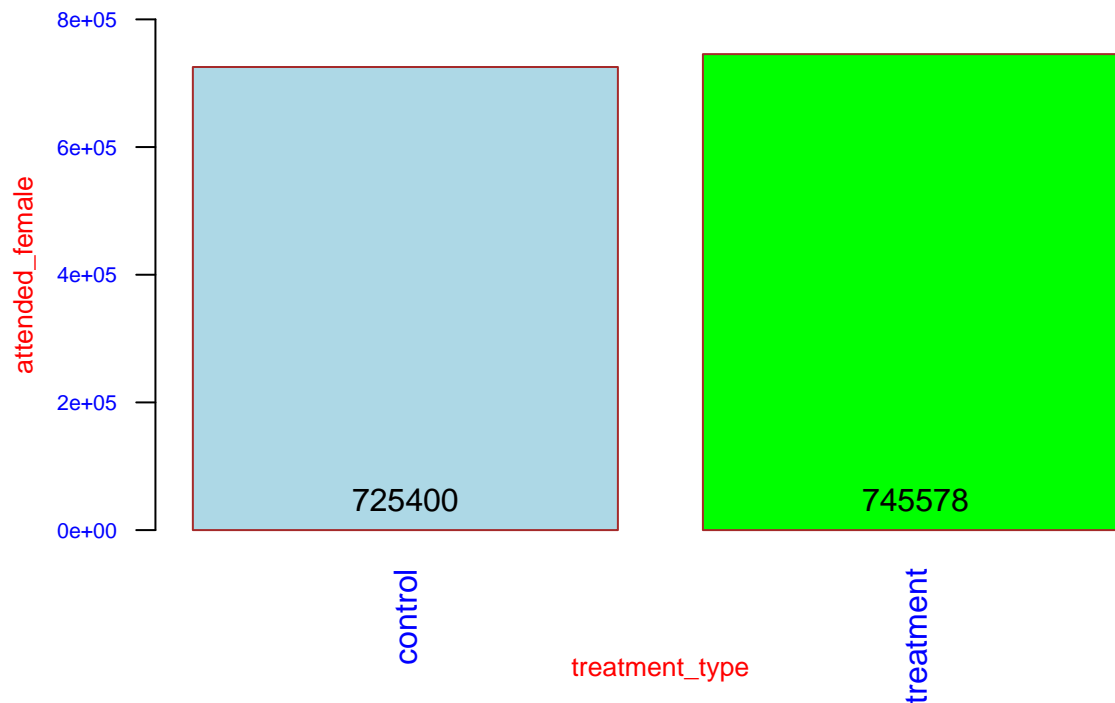
```
which(!complete.cases(ds_sdmerger1))
```

```
## integer(0)
```

```
#Create a well labelled graphs showing the difference in attendance between treatment  
#and control schools.
```


```
bp1 <- barplot(tapply(ds_sdmerger1$attended_female, ds_sdmerger1$treatment_type, FUN=sum),
               xlab="treatment_type",ylab="attended_female",col=c("lightblue", "green"),
               main="Sum Attended Female Barplot chart",border="brown" , col.axis = 'blue', col.lab = 'red',
               cex.axis = 0.7, cex.lab = 0.8, las = 2, ylim=c(0, 800000))
text(bp1, 0, tapply(ds_sdmerger1$attended_female, ds_sdmerger1$treatment_type, FUN=sum),
     cex=1,pos=3)
```

Sum Attended Female Barplot chart




```
bp2 <- barplot(tapply(ds_sdmerger1$attended_male, ds_sdmerger1$treatment_type, FUN=sum),
  xlab="treatment_type",ylab="attended_male",col=c("yellow", "orange"),
  main="Sum Attended Male Barplot chart",border="brown" , col.axis = 'blue', col.lab = 'red',
  cex.axis = 0.7, cex.lab = 0.8, las = 2, ylim=c(0, 1000000))
text(bp2, 0, tapply(ds_sdmerger1$attended_male, ds_sdmerger1$treatment_type, FUN=sum),
  cex=1,pos=3)
```

Sum Attended Male Barplot chart




```
bp3 <- barplot(tapply(ds_sdmerger1$total_attended, ds_sdmerger1$treatment_type, FUN=sum),
  xlab="treatment_type",ylab="total_attended",col=c("grey", "purple"),
  main="Total Sum Attended Barplot chart",border="brown" , col.axis = 'blue', col.lab = 'red',
  cex.axis = 0.7, cex.lab = 0.8, las = 2, ylim=c(0, 2000000))
text(bp3, 0, tapply(ds_sdmerger1$total_attended, ds_sdmerger1$treatment_type, FUN=sum),
  cex=1,pos=3)
```


Total Sum Attended Barplot chart




```
#relationship between treatment_type and total_attended by means of a boxplot  
boxplot(attended_female ~ treatment_type,  
col=c("lightblue", "green"),ds_sdmerger1,  
col.axis = 'blue', col.lab = 'red', border="brown",  
main="Sum Attended Female BoxPlot")
```




```
boxplot(attended_male ~ treatment_type,  
col=c("yellow", "orange"),ds_sdmerger1,  
col.axis = 'blue', col.lab = 'red', border="brown",  
main="Sum Attended Male BoxPlot")
```

Sum Attended Male BoxPlot



```
boxplot(total_attended ~ treatment_type,  
col=c("gray","purple"),ds_sdmerger1,  
col.axis = 'blue', col.lab = 'red', border="brown",  
main="Total Sum Attended BoxPlot")
```


Total Sum Attended BoxPlot



```
#finding a fitting distribution for the total attended variable
library(car)
library(MASS) #So that distributions that must be non-zero can make sense of my data

qqp(ds_sdmerger1$total_attended+1, "norm", main="Normal model")
```


Normal model



```
## [1] 6180 2573
```

```
qqp(ds_sdmerger1$total_attended+1, "lnorm", main="LogNormal model") #lnorm means lognormal
```

LogNormal model



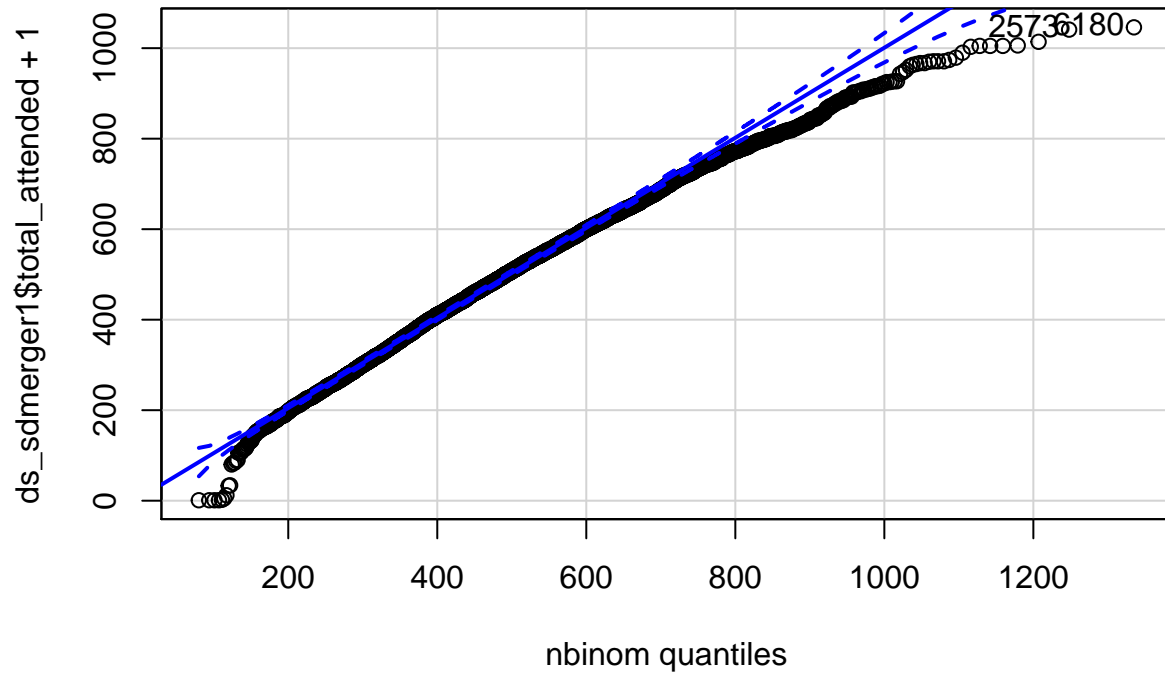
```
## [1] 6180 2573
```

```
# qqp requires estimates of the parameters of the negative binomial, Poisson  
# and gamma distributions. You can generate estimates using the fitdistr  
# function.
```

```
#negative binomial and gamma distributions can only handle positive numbers.  
#Poisson distribution can only handle positive whole numbers.  
#Binomial and Poisson distributions are different from the others because they  
#are discrete rather than continuous, which means they quantify distinct,  
#countable events or the probability of these events
```

```
nbinom <- fitdistr(ds_sdmerger1$total_attended+1, "Negative Binomial")  
qqp(ds_sdmerger1$total_attended+1, "nbinom", size = nbinom$estimate[[1]], mu =  
    nbinom$estimate[[2]], main="Negative Binomial model")
```

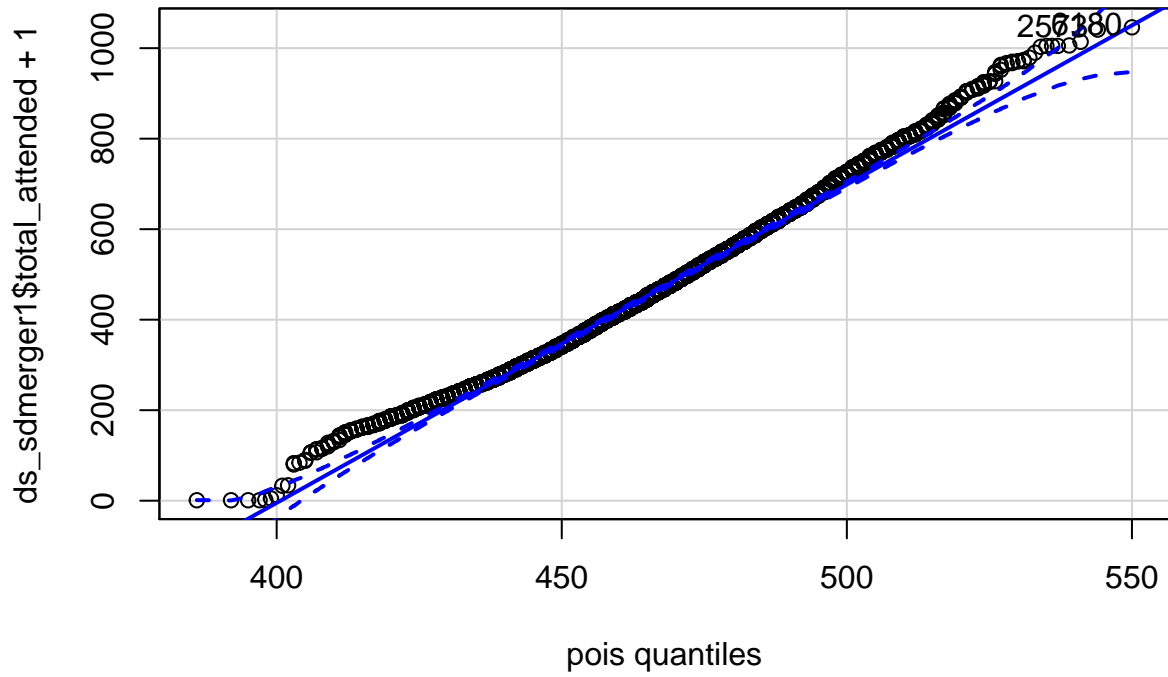
Negative Binomial model



```
## [1] 6180 2573
```

```
pois <- fitdistr(ds_sdmerger1$total_attended+1, "Poisson")  
qqp(ds_sdmerger1$total_attended+1, "pois", lambda=pois$estimate, main="Poisson model")
```


Poisson model



```
## [1] 6180 2573
```

```
gamma <- fitdistr(ds_sdmerger1$total_attended+1, "gamma")
qqp(ds_sdmerger1$total_attended+1, "gamma", shape = gamma$estimate[[1]], rate =
    gamma$estimate[[2]], main="Gamma model")
```


Gamma model



```
## [1] 6180 2573
```

```
#Armed with the knowledge of which probability distribution fits best,  
#we can try fitting a model
```

```
#A mixed model is similar in many ways to a linear model.  
#It estimates the effects of one or more explanatory variables on a response variable.  
#The output of a mixed model will give you a list of explanatory values, estimates and  
#confidence intervals of their effect sizes, p-values for each effect, and  
#at least one measure of how well the model fits.  
#You should use a mixed model instead of a simple linear model when you have a variable  
#that describes your data sample as a subset of the data you could have collected.  
  
#In a mixed model, we add one or more random effects to our fixed effects.  
#These random effects essentially give structure to the error term.  
#this characterizes idiosyncratic variation that is due to individual differences.  
  
#linear models are "fixed-effects-only" models. They have one or more fixed effects and a  
#general error term.  
  
#If data is normally distributed, we can use a linear mixed model (LMM).  
#load the lme4 package and make a call to the function lmer.  
#The first argument to the function is a formula that takes the form y ~ x1 + x2 ...etc.,  
#where y is the response variable and x1, x2, etc. are explanatory variables.  
#Random effects are added in with the explanatory variables.  
#Crossed random effects take the form (1 | r1) + (1 | r2) ...
```

```
#while nested random effects take the form (1 | r1 / r2).
```

*#The next argument is where you designate the data frame your variables come from.
 #This is where you can designate whether the mixed model will estimate the parameters
 #using maximum likelihood or restricted maximum likelihood.
 #If your random effects are nested, or you have only one random effect, and if your data
 #are balanced (i.e., similar sample sizes in each factor group) set REML to FALSE,
 #because you can use maximum likelihood.
 #If your random effects are crossed, don't set the REML argument because it defaults
 #to TRUE anyway.*

```
library(lme4)
#1. We construct the null model first.
#H0 (called the null hypothesis): There is no relationship between the two variables.

lmm_treatment.null <- lmer(total_attended ~ total_enrolled + (1 | district_id),
                           data = ds_sdmerger1, REML = FALSE)
summary(lmm_treatment.null)
```

```
## Linear mixed model fit by maximum likelihood ['lmerMod']
## Formula: total_attended ~ total_enrolled + (1 | district_id)
## Data: ds_sdmerger1
##
##      AIC      BIC    logLik deviance df.resid
## 88998.4 89025.8 -44495.2 88990.4      6944
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -3.5915 -0.7191 -0.0497  0.6586  3.6938
##
## Random effects:
## Groups      Name      Variance Std.Dev.
## district_id (Intercept) 1797      42.39
## Residual                21121    145.33
## Number of obs: 6948, groups: district_id, 27
##
## Fixed effects:
##              Estimate Std. Error t value
## (Intercept)  4.103e+02  9.149e+00  44.84
## total_enrolled 5.773e-02  3.825e-03  15.09
##
## Correlation of Fixed Effects:
##              (Intr)
## total_nrlld -0.410
```

*#The Anova function does a Wald test, which tells us how confident we are of our
 #estimate of the effect of total enrolled on total attended, and the p-value tells
 #me that I should not be confident at all.*

```
Anova(lmm_treatment.null)
```

```
## Analysis of Deviance Table (Type II Wald chisquare tests)
##
## Response: total_attended
##              Chisq Df Pr(>Chisq)
## total_enrolled 227.82  1 < 2.2e-16 ***
```

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

#2.we construct the full model next
#H1 (called the alternative hypothesis): There exist a relationship between the two variables.

lmmtreatment <- lmer(total_attended ~ treatment + total_enrolled + (1 | district_id),
                     data = ds_sdmerger1, REML = FALSE)
summary(lmmtreatment)

## Linear mixed model fit by maximum likelihood ['lmerMod']
## Formula: total_attended ~ treatment + total_enrolled + (1 | district_id)
## Data: ds_sdmerger1
##
##      AIC      BIC   logLik deviance df.resid
## 88994.6 89028.8 -44492.3 88984.6      6943
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -3.6227 -0.7186 -0.0535  0.6572  3.6659
##
## Random effects:
## Groups      Name      Variance Std.Dev.
## district_id (Intercept) 1797      42.39
## Residual                21103    145.27
## Number of obs: 6948, groups: district_id, 27
##
## Fixed effects:
##              Estimate Std. Error t value
## (Intercept)  4.060e+02  9.320e+00  43.558
## treatment    8.406e+00  3.486e+00   2.411
## total_enrolled 5.785e-02  3.823e-03  15.131
##
## Correlation of Fixed Effects:
##              (Intr) trtmnt
## treatment    -0.191
## total_nrlld -0.405  0.014

#The Anova function tells us how confident we are of our estimate of the effect of
#treatment and total enrolled on total attended, and the p-value tells
#me that I should not be confident at all.
Anova(lmmtreatment)

## Analysis of Deviance Table (Type II Wald chisquare tests)
##
## Response: total_attended
##              Chisq Df Pr(>Chisq)
## treatment      5.814  1    0.0159 *
## total_enrolled 228.955  1    <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

lmmtreat <- lmer(total_attended ~ treatment + (1 | district_id),
                 data = ds_sdmerger1, REML = FALSE)
summary(lmmtreat)
```

```
## Linear mixed model fit by maximum likelihood ['lmerMod']
## Formula: total_attended ~ treatment + (1 | district_id)
## Data: ds_sdmerger1
##
##      AIC      BIC    logLik deviance df.resid
## 89217.8 89245.2 -44604.9 89209.8     6944
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -3.2597 -0.7178 -0.0352  0.6498  3.6482
##
## Random effects:
## Groups      Name      Variance Std.Dev.
## district_id (Intercept) 1905     43.65
## Residual              21797    147.64
## Number of obs: 6948, groups: district_id, 27
##
## Fixed effects:
##              Estimate Std. Error t value
## (Intercept)  463.090      8.766  52.830
## treatment      7.691      3.542   2.171
##
## Correlation of Fixed Effects:
##              (Intr)
## treatment -0.201
```

```
#The Anova function tells us how confident we are of our estimate of the effect of
#treatment on total attended, and the p-value tells
#me that I should not be confident at all.
Anova(lmmtreat)
```

```
## Analysis of Deviance Table (Type II Wald chisquare tests)
##
## Response: total_attended
##           Chisq Df Pr(>Chisq)
## treatment 4.7136 1    0.02993 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
#The coefficient "treatment" is the slope for the categorical effect of providing meals.
#7.691 means that to go from "control" to "treatment", total attendance increases
#by around 8 pupils.
#attendance is lower in control(no meals provided) than in treatment(meals provided),
#by about 8 pupils.
```

```
#oftentimes, model intercepts are not particularly meaningful.
#But this Model intercept is 463.090. If you look back at the boxplot that we constructed
#earlier, you can see that the value 463.090 seems to fall halfway between control
#and treatment(and this is indeed what this intercept represents).
#It's the average of our data for the informal condition
```

```
#compared to the other model with the fixed effect total enrolled added, the intercept
#is particularly off as we didn't inform our model that there's total enrolled in
#our dataset. the intercept reduces to 406. The coefficient for the effect of treatment
#increased to 8.406 from 7.691
```

```

#If you want to interpret these results, you'll most likely need to report some kind
#of p-value.
#P-value for treatment in the models are significant(less than 0.05) and reduces
#from 0.02993 to 0.0159.
#Unfortunately, p-values for mixed models aren't as straightforward as they are
#for the linear model.

#Rather than getting a p-value straightforwardly from your model, we get a p-value from a
#comparison of two models.

#Thus we focus on the Likelihood Ratio Test as a means to attain p-values.
#Likelihood is the probability of seeing the data you collected given your model.
#The logic of the likelihood ratio test is to compare the likelihood of two models
#with each other. First, the model without the factor that you're interested in (the
#null model), then the model with the factor that you're interested in.

```

```

#3. We have two models to compare with each other - one with the effect in
#question, one without the effect in question.
#We perform the likelihood ratio test using the anova() function:

```

```
anova(lmmtreatment.null,lmmtreatment)
```

```

## Data: ds_sdmerger1
## Models:
## lmmtreatment.null: total_attended ~ total_enrolled + (1 | district_id)
## lmmtreatment: total_attended ~ treatment + total_enrolled + (1 | district_id)
##           Df    AIC    BIC logLik deviance Chisq Chi Df Pr(>Chisq)
## lmmtreatment.null  4 88998 89026 -44495     88990
## lmmtreatment      5 88995 89029 -44492     88985 5.8116      1 0.01592 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

```

#we compared a full model (with the fixed effects in question) against a reduced(null)
#model without the effects in question.
#we conclude that a fixed effect treatment is significant as the difference between the
#likelihood of these two models is significant.

```

```

#treatment(providing primary school pupils with a free meal on school days) affected total
#attendance (chisq.(1)=5.8116, p=0.01592), increasing it by about 8.406(8pupils) ± 3.486(3 pupils)

```

```
coef(lmmtreatment)
```

```

## $district_id
##      (Intercept) treatment total_enrolled
## 3      345.5238  8.405591      0.0578538
## 5      373.2582  8.405591      0.0578538
## 6      391.9255  8.405591      0.0578538
## 16     375.2961  8.405591      0.0578538
## 17     445.5646  8.405591      0.0578538
## 22     427.8311  8.405591      0.0578538
## 39     380.9146  8.405591      0.0578538
## 42     353.3707  8.405591      0.0578538
## 43     449.6900  8.405591      0.0578538
## 44     410.6224  8.405591      0.0578538
## 47     507.9889  8.405591      0.0578538

```

```
## 50      449.4153  8.405591      0.0578538
## 57      429.5080  8.405591      0.0578538
## 58      417.6188  8.405591      0.0578538
## 62      424.2424  8.405591      0.0578538
## 65      382.6116  8.405591      0.0578538
## 67      396.6649  8.405591      0.0578538
## 68      353.9196  8.405591      0.0578538
## 77      405.3434  8.405591      0.0578538
## 78      340.3056  8.405591      0.0578538
## 79      410.7082  8.405591      0.0578538
## 80      378.1203  8.405591      0.0578538
## 93      464.7975  8.405591      0.0578538
## 97      446.7653  8.405591      0.0578538
## 104     463.5627  8.405591      0.0578538
## 107     369.6693  8.405591      0.0578538
## 108     365.8604  8.405591      0.0578538
##
## attr("class")
## [1] "coef.mer"
```

#You see that each district is assigned a different intercept given that we've told the #model with "(1|district_id)"

#Note also that the fixed effects (treatment and total_enrolled) are all the same for all #district_id. Our model is what is called a random intercept model.

#In this model, we account for baseline-differences in total attended, but we assume that #whatever the effect of treatment is, it's going to be the same for all subjects and items.

#But is that a valid assumption? often times it's not - it is quite expected #that some district-id would elicit more or less treatments. That is, the effect of #providing meals might be different for different district-id.

#Thus what we need is a random slope model, where district_id is not only allowed to have #differing intercepts, but where they are also allowed to have different slopes for the #effect of providing meals

```
lmmtreatmentRS <- lmer(total_attended ~ treatment + total_enrolled +
  (1+treatment|district_id), data = ds_sdmerger1, REML = FALSE)
summary(lmmtreatmentRS)
```

```
## Linear mixed model fit by maximum likelihood ['lmerMod']
## Formula: total_attended ~ treatment + total_enrolled + (1 + treatment |
##   district_id)
## Data: ds_sdmerger1
##
##      AIC      BIC    logLik deviance df.resid
## 88977.1 89025.0 -44481.5  88963.1      6941
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -3.5108 -0.7124 -0.0544  0.6558  3.7093
##
## Random effects:
## Groups      Name             Variance Std.Dev. Corr
## district_id (Intercept) 1643.7   40.54
```

```
##           treatment      578.4   24.05   0.01
## Residual                20959.0  144.77
## Number of obs: 6948, groups:  district_id, 27
##
## Fixed effects:
##           Estimate Std. Error t value
## (Intercept)    4.061e+02  9.007e+00  45.084
## treatment      9.052e+00  5.794e+00   1.562
## total_enrolled 5.743e-02  3.816e-03  15.049
##
## Correlation of Fixed Effects:
##           (Intr) trtmnt
## treatment    -0.111
## total_nrllld -0.418  0.008
```

```
lmmtreatmentRS.null <- lmer(total_attended ~ total_enrolled + (1+treatment|district_id),
  data = ds_sdmerger1, REML = FALSE)
summary(lmmtreatmentRS.null)
```

```
## Linear mixed model fit by maximum likelihood ['lmerMod']
## Formula: total_attended ~ total_enrolled + (1 + treatment | district_id)
## Data: ds_sdmerger1
##
##      AIC      BIC    logLik deviance df.resid
## 88977.4 89018.5 -44482.7 88965.4      6942
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -3.4941 -0.7123 -0.0532  0.6551  3.7198
##
## Random effects:
## Groups      Name      Variance Std.Dev. Corr
## district_id (Intercept) 1651      40.63
##              treatment    656      25.61  -0.01
## Residual                20959      144.77
## Number of obs: 6948, groups:  district_id, 27
##
## Fixed effects:
##           Estimate Std. Error t value
## (Intercept)    4.077e+02  8.959e+00  45.51
## total_enrolled 5.736e-02  3.816e-03  15.03
##
## Correlation of Fixed Effects:
##           (Intr)
## total_nrllld -0.420
```

#The notation "(1+treatment | district_id)" means that you tell the model to expect differing #baseline-levels of total_attended (the intercept, represented by 1) as well as differing #responses to the main factor which is "treatment" in this case

```
anova(lmmtreatmentRS.null,lmmtreatmentRS)
```

```
## Data: ds_sdmerger1
## Models:
## lmmtreatmentRS.null: total_attended ~ total_enrolled + (1 + treatment | district_id)
```

```
## lmmtreatmentRS: total_attended ~ treatment + total_enrolled + (1 + treatment |
## lmmtreatmentRS:      district_id)
##           Df    AIC    BIC logLik deviance  Chisq Chi Df Pr(>Chisq)
## lmmtreatmentRS.null  6 88977 89019 -44483    88965
## lmmtreatmentRS      7 88977 89025 -44482    88963 2.3481      1    0.1254
```

*#we conclude that a fixed effect treatment for random slope model is not significant as the
#difference between the likelihood of these two models is 0.1254.(greater than 0.05)*

```
coef(lmmtreatmentRS)
```

```
## $district_id
##      (Intercept)      treatment total_enrolled
## 3      340.7621    19.9021088      0.05742947
## 5      374.0672     7.7820314      0.05742947
## 6      380.0712    34.5099362      0.05742947
## 16     384.1282   -9.0841655      0.05742947
## 17     447.6537     4.4681956      0.05742947
## 22     416.4717    33.4256029      0.05742947
## 39     380.9118     9.4537385      0.05742947
## 42     356.5562     2.9167096      0.05742947
## 43     460.0651  -13.1021611      0.05742947
## 44     424.7478  -20.6504402      0.05742947
## 47     501.1083    22.9982718      0.05742947
## 50     441.6161    25.4786518      0.05742947
## 57     421.2400    26.3587508      0.05742947
## 58     433.7818  -25.0150803      0.05742947
## 62     414.4051    29.8114630      0.05742947
## 65     383.3124     7.9127314      0.05742947
## 67     400.9756     0.2866173      0.05742947
## 68     357.2551     2.6064692      0.05742947
## 77     403.1450    13.8545405      0.05742947
## 78     356.9068  -25.3765606      0.05742947
## 79     422.9093  -16.4483004      0.05742947
## 80     386.7669    -8.6627213      0.05742947
## 93     457.3294    24.7576474      0.05742947
## 97     427.9750    49.0141645      0.05742947
## 104     457.9225    20.7703352      0.05742947
## 107     361.6458    26.5974243      0.05742947
## 108     370.4186   -0.1722125      0.05742947
##
## attr(,"class")
## [1] "coef.mer"
```